

Zhongqi Yue

INTERESTS: LLM POST-TRAINING, MULTIMODAL LLM, CAUSAL GENERALIZATION

Computer Science and Engineering, Chalmers University of Technology

yuez002@gmail.com | [Home Page](#) | [yue-zhongqi](#)

Research Highlights

My research lies at the intersection of *post-training for LLMs*, *multimodal LLMs*, and *causal generalization*.

- My most recent project [EARL](#) proposes a new paradigm for post-training LLMs with external environments, which decouples environment interactions from language by internalizing them in an expanded action space.
- I proposed a diffusion-based unsupervised representation learning method (ICLR 2024) and subsequently pioneered the [Selftok](#) image tokenizer (**CVPR 2025 Best Student Paper Honorable Mention**). I supervised multimodal data generation/selection (multiple **oral/highlight** papers), and directed large-scale multimodal LLMs pre-training and post-training on 150+ GPUs ([technical report](#)).
- Previously, I studied disentangled representation learning through contrastive objectives, and explored machine generalization problems (e.g., zero-/few-shot learning, domain adaptation) from a causal perspective.

Professional Experience

Wallenberg-NTU Presidential Postdoctoral Fellowship

Singapore & Sweden

PRESIDENTIAL POSTDOCTORAL FELLOW, PRINCIPAL INVESTIGATOR

01.2024-02.2026 (Expected)

- 01.2024 - 01.2025: MReal Lab, Nanyang Technological University, advised by Prof. Hanwang Zhang;
- 02.2025 - 02.2026: Healthy AI Lab, Chalmers University of Technology, advised by Assoc Prof. Fredrik Johansson.

Sea AI Lab (SAIL)

Singapore

RESEARCH INTERN

10.2023 - 12.2023

- Supervisor: Asst Prof. Pan Zhou

Alibaba-NTU Singapore Joint Research Institute

Singapore

ALIBABA PHD STUDENT

08.2019 - 08.2023

- Advisor: Xian-Sheng Hua (IEEE Fellow)

Nanyang Technological University

Singapore

PH.D IN COMPUTER SCIENCE

08.2019 - 12.2023

- Advisor: Prof. Hanwang Zhang
- Thesis: Causal View of Generalization

Selected Publications

[**NeurIPS 2025 Math-AI**] Zhongqi Yue*, Weishi Wang*, Yundaichuan Zhan, Juncheng Li, Daniel Dahlmeier, Fredrik D. Johansson. Expanding the Action Space of LLMs to Reason Beyond Language. 2025.

[**CVPR 2025**] Kaihang Pan*, Wang Lin*, Zhongqi Yue*, Tenglong Ao, Liyu Jia, Wei Zhao, Juncheng Li, Siliang Tang, Hanwang Zhang. Generative Multimodal Pretraining with Discrete Diffusion Timestep Tokens. (**Best Student Paper Honorable Mention** 7/13008).

[**Technical Report**] Bohan Wang, Zhongqi Yue, Fengda Zhang, Shuo Chen, Li'an Bi, Junzhe Zhang, Xue Song, Kennard Yanting Chan, Jiachun Pan, Weijia Wu, Mingze Zhou, Wang Lin, Kaihang Pan, Saining Zhang, Liyu Jia, Wentao Hu, Wei Zhao, Hanwang Zhang. Selftok: Discrete visual tokens of autoregression, by diffusion, and for reasoning. 2025.

[**AAAI 2026**] Xue Song, Zhongqi Yue, Jiequan Cui, Hanwang Zhang, Jingjing Chen, Yu-Gang Jiang. Object Fusion via Diffusion Time-step for Customized Image Editing with Single Example. (**Oral Presentation**).

[**EurIPS 2025 MMRL4H & MedEurIPS**] Herman Bergström, Zhongqi Yue, Fredrik D. Johansson. When are radiology reports useful for training medical image classifiers? (**Oral Presentations at Both Workshops**)

[**NeurIPS 2025**] Bohan Wang, Mingze Zhou, Zhongqi Yue, Wang Lin, Kaihang Pan, Liyu Jia, Wentao Hu, Wei Zhao, Hanwang Zhang. Selftok-Zero: Reinforcement Learning for Visual Generation via Discrete and Autoregressive Visual Tokens.

[ICCV 2025] Qifan Yu*, Zhebei Shen*, **Zhongqi Yue***, Yang Wu, Wenqiao Zhang, Yunfei Li, Juncheng Li, Siliang Tang, Yueling Zhuang. Mastering Collaborative Multi-modal Data Selection: A Focus on Informativeness, Uniqueness, and Representativeness. (**Highlight 262/11239**).

[CVPR 2025] Qifan Yu*, Wei Chow*, **Zhongqi Yue***, Kaihang Pan, Yang Wu, Xiaoyang Wan, Juncheng Li, Siliang Tang, Hanwang Zhang, Yueling Zhuang. AnyEdit: Mastering unified high-quality image editing for any idea. (**Oral Presentation 96/13008**).

[ICLR 2024] **Zhongqi Yue**, Jiankun Wang, Qianru Sun, Lei Ji, Eric I-Chao Chang, Hanwang Zhang. Exploring Diffusion Time-steps for Unsupervised Representation Learning.

[NeurIPS 2021] Tan Wang, **Zhongqi Yue**, Jianqiang Huang, Qianru Sun, Hanwang Zhang. Self-Supervised Learning Disentangled Group Representation as Feature. (**Spotlight Presentation 260/9122; 2022 PREMIA Best Paper**)

[ICCV 2021] **Zhongqi Yue**, Qianru Sun, Xian-Sheng Hua, Hanwang Zhang. Transporting Causal Mechanisms for Unsupervised Domain Adaptation. (**Oral Presentation 210/6236**)

Other Publications

[NeurIPS 2025] Minghe Gao*, **Zhongqi Yue***, Wenjie Yan, Yihao Hu, Wei Ji, Siliang Tang, Jun Xiao, Tat-Seng Chua, Yueling Zhuang, Juncheng Li. Counterfactual Evolution of Multimodal Datasets via Visual Programming.

[ICCV 2025] Minghe Gao, Xuqi Liu, **Zhongqi Yue**, Yang Wu, Shuang Chen, Juncheng Li, Siliang Tang, Fei Wu, Tat-Seng Chua, Yueling Zhuang. Benchmarking multimodal cot reward model stepwise by visual program.

[CVPR 2024] **Zhongqi Yue**, Pan Zhou, Richang Hong, Hanwang Zhang, Qianru Sun. Few-shot Learner Parameterization by Diffusion Time-steps.

[NeurIPS 2023] **Zhongqi Yue**, Qianru Sun, Xian-Sheng Hua, Hanwang Zhang. Make the U in UDA Matter: Invariant Consistency Learning for Unsupervised Domain Adaptation. (**WILDS 2.0 Leaderboard 1st Place**).

[ICCV 2023] Jiali Ma, **Zhongqi Yue**, Tomoyuki Kagaya, Tomoki Suzuki, Karlekar Jayashree, Sugiri Pranata, Hanwang Zhang. Invariant Feature Regularization for Fair Face Recognition.

[ICCV 2023] Yanghao Wang, **Zhongqi Yue**, Xian-Sheng Hua, Hanwang Zhang.
Random Boxes Are Open-world Object Detectors.

[CVPR 2023] Hui Lv, **Zhongqi Yue**, Qianru Sun, Bin Luo, Zhen Cui, Hanwang Zhang. Unbiased Multiple Instance Learning for Weakly Supervised Video Anomaly Detection.

[CVPR 2021] **Zhongqi Yue***, Tan Wang*, Qianru Sun, Xian-Sheng Hua, Hanwang Zhang. Counterfactual Zero-Shot and Open-Set Visual Recognition.

[NeurIPS 2020] **Zhongqi Yue**, Hanwang Zhang, Qianru Sun, Xian-Sheng Hua. Interventional Few-Shot Learning.

Wang Lin, Liyu Jia, Wentao Hu, Kaihang Pan, **Zhongqi Yue**, Wei Zhao, Jingyuan Chen, Fei Wu, Hanwang Zhang. Reasoning physical video generation with diffusion timestep tokens via reinforcement learning. 2025.

Yuxuan Liu, **Zhongqi Yue**, Kun Kuang, Fengda Zhang, Hanwang Zhang. Open-Set Recognition Beyond Vision-Language Pre-training. 2025.

Academic Services

INVITED TALKS

2025. *Generative Multimodal Pretraining with Discrete Diffusion Timestep Tokens*. Oral presentation at CVPR.

2025. *AnyEdit: Mastering unified high-quality image editing for any idea..* Oral presentation at CVPR.

2025. *Large Language Models: How Machines Learned to Talk and Think*. Invited talk at Huawei Singapore Research Center.

2025. *The Rise of Multimodal LLM: How Machines Learned to See and Draw*. Invited talk at Huawei Singapore Research Center.

2024. *Tokenizer is Key to Next-Gen AR Modelling*. Invited talk at Huawei Singapore Research Center.

2024. *Visual Autoregressive Modelling*. Invited talk by NUS NExT Research Centre.

2024. *Visual Autoregressive Modelling*. Invited talk by Skywork AI.

2021. *Interventional Learning*. Invited talk by the Lab for Media Search (LMS) in NUS and NExT Research Centre.

2020. *Interventional Few-Shot Learning*. Invited talk by Jiantao Jiao's Lab in UC Berkeley.

TEACHING EXPERIENCE

- Sem 1, AY2020-21 **Introduction to Computational Thinking and Programming**, Teaching Assistant
- Sem 1, AY2021-22 **Microprocessor System Design and Development**, Teaching Assistant
- Sem 2, AY2022-23 **Artificial Intelligence**, Teaching Assistant

PEER REVIEW

Conference & journal reviewer for NeurIPS, CVPR, ICCV, ICML, UAI, ECCV, IJCAI, ACM MM, TPAMI, TMLR, IJCV, T-RO.

Awards & Scholarships

- 2019-2023 **Alibaba Talent Program Candidate**, Alibaba Cloud
- 2015-2016 **NTU President Research Scholar**, Nanyang Technological University
- 2011-2017 **MOE SM2 Scholoarship** (200 candidates annually in China), MOE Singapore
- 2017 **Shape The Future Global Innovation Winner** (2/300), Tata Communications
- 2017 **AWS Hack Day 2017 2nd Prize**, Amazon Web Services
- 2015-2017 **Dean's List and Academic Excellence Award**, Nanyang Technological University
- 2018 **Certified SCRUM Master**, SCRUM Alliance
- 2018 **AWS Certified Solution Architect & Developer**, Amazon Web Services

Outreach & Professional Development

- 2017 **Guest performer**, NTU Guitar Ensemble Annual Concert
- 2016 **Organizer**, Oversea Visit to NTU and NTNU in Taiwan
- 2015 **President**, NTU Cultural Activities Club, Harmonica Band
- 2013 **Team Leader**, Volunteering for Singapore School for the Deaf

Skills

Recently Used: Python, PyTorch, Tensorflow, L^AT_EX;

Previously Used: C++, Java, Javascript, Frontend & backend development.